

Modernizing the Emergent Money Simulation – Comprehensive Plan

Introduction

The *Emergent Money* simulation is a legacy agent-based economic model originally implemented in C (file `adamsmithfeb17B.c`, based on the 2010 report *Emergent Money* ¹). It simulates a barter economy with *emergent* currency-like behavior among a population of agents. The goal is to modernize this simulation using contemporary tools – specifically leveraging AI-assisted coding (OpenAI Codex), rewriting logic in Python for maintainability, harnessing parallelism/GPU computing for performance, and building an interactive dashboard-style user interface (UI). This report outlines a comprehensive plan for the modernization effort, including an assessment of the legacy code's complexity, the proposed technology stack (GPU libraries, agent-based frameworks, UI frameworks), the feasibility of a Python re-implementation, and a stepwise development strategy. It also covers integration of the simulation kernel with visualization, future web deployment considerations, testing/validation against the original model's results, and a source control strategy.

Assessment of the Legacy Simulation Code

The existing simulation (`adamsmithfeb17B.c`) comprises roughly 2,600 lines of C code, structured around global data structures that represent a “market” and a fixed-size array of agents. Key characteristics of the legacy implementation include:

- **Integer-Heavy Arithmetic:** Most calculations use integer or fixed-point math for performance. The code scales many values by 10 (see `INTMULTIPLIER` in the code) and rounds off small deviations ². For example, basic needs are defined as `i * i * INTMULTIPLIER` in the initialization, using integers to represent scaled quantities. This approach improved speed on older hardware (avoiding costly floating-point operations), but it reduces clarity and precision. The project report explicitly notes that “the software is written in C using integer math in most algorithms” and that many numbers are “ten-folded” for computation ².
- **Global Memory Architecture:** All agents’ data is stored in large global arrays and structs, enabling direct memory access. For instance, a global `MARKETSTRUCTURE market` holds economy-wide aggregates (production totals, prices, etc.), and an array `individual[POPULATION]` holds each agent’s state. Each agent structure contains numerous fields – e.g. arrays for needs, surplus, stock limits, prices for each of 31 goods (“SKILLS”), social network relations, and so on ³ ⁴. This design allows fast memory access and in-place updates (agents share memory for interactions), but it tightly couples logic to specific data layouts and makes the code monolithic and harder to modify.
- **Complex Logic with Multiple Interacting Components:** The simulation covers agent consumption, production (normal, surplus, and leisure production rounds), barter negotiations between socially connected agents, price formation, and inventory spoilage. The code is divided into many functions (e.g. `initmarket`, `initpopulation`, `getfriendexchangeindexes`, `makesurplusdeals`, `produceneed`, `consumesurplus`, `adjustpurchaseprice`,

endmyperiod, etc.), which are orchestrated in a main cycle loop (referred to as “periods” or cycles) ⁵. The *Emergent Money* report describes that each simulation cycle updates almost the entire “database” of agent states ⁶. This means that each period, every agent recalculates needs, produces goods, engages in trades with friends, updates prices, and so forth, leading to heavy computation. The code uses nested loops over agents, their social connections (up to 10 friends per agent), and 31 possible goods, resulting in many iterations per cycle. For example, trading logic examines potential exchanges of each good with each friend, leading to nested loops over goods and friend links ⁷ ⁸.

- **Performance Characteristics:** With the chosen integer math and in-memory data, the original simulation was optimized for speed on a single CPU core. The report notes that even with these optimizations, the simulation is time-consuming: a test run with ~3000 agents (each having a network of friends and multiple needs) takes a few hours on a “regular PC” ⁹. (The default configuration in code uses 300 agents, likely for quicker tests.) This underscores that the algorithms are computationally intensive, though the agent count and network size were relatively limited. The code’s use of fixed-size arrays and global state implies it was not trivial to parallelize in its original form – it was essentially single-threaded, relying on CPU speed and memory efficiency.
- **Output and Monitoring:** The legacy program outputs aggregated statistics each cycle (to a file and/or console). For example, it prints metrics like spoils, transaction costs, production, etc., per period ¹⁰. It also tracks a “test individual” (one agent) to output detailed behavior for analysis ¹¹. This indicates that while running, the simulation provided some observability, but it lacked a rich interactive interface – analysis was done by studying log files or console prints post hoc. Indeed, the authors spent “hundreds of hours” examining agent behaviors through such outputs ⁹.

Summary of Complexity: Overall, the legacy code is efficient but *dense*. It intermixes economic logic with low-level details (array indices, manual accumulation of statistics, etc.). The heavy use of integers and global mutable state improved performance but makes the code less transparent. Modernizing this will involve carefully disentangling and re-implementing logic in a high-level manner, while ensuring we preserve the emergent economic behaviors (e.g. growth, cycles, barter dynamics) described in the original study ¹². The challenge is to maintain correctness and performance *and* add flexibility (for new UI controls and possibly larger scale).

Modernization Goals and Challenges

Based on the project’s objectives, the modernization effort has several key goals:

- **Full Code Rewrite with AI Assistance:** Reimplement the simulation from scratch in a modern language (preferably Python). The plan is to leverage OpenAI Codex (or similar GPT-based coding assistants) to generate the bulk of the new code. The user (and AI assistant, e.g. ChatGPT) will guide the architecture and ensure the generated code aligns with the intended design. The idea is to *pair-program* with AI: using it to accelerate writing boilerplate and even complex logic, while the human ensures coherence and accuracy. This approach can significantly speed up development, but it requires a clear plan for module separation and incremental integration (to avoid getting lost in one giant AI-generated code dump).
- **Shift to Python for Maintainability:** Python is chosen for its readability and vast ecosystem. The new code should be more maintainable, modular, and easier to extend than the C code.

However, a direct concern is performance: Python is much slower in raw loop execution than C. We must address this by using optimized libraries, vectorized operations, and possibly just-in-time compilation or GPU offloading. The *feasibility* of using Python will depend on whether we can achieve acceptable performance for the agent population sizes of interest, which might range from a few hundred to many thousands. We assume that by leveraging parallel hardware (multi-core CPUs and GPUs), we can compensate for Python's overhead.

- **Higher Precision and Clarity:** In moving away from “integer math,” we can use Python's float (double precision) or even decimal types for more natural calculations. This will simplify the logic (no more manual $\times 10$ scaling) and improve numeric fidelity (avoiding the rounding off of “small deviations” mentioned in the report ¹³). It's important to double-check that using floating-point doesn't introduce significant divergence in simulation outcomes – we will validate that the new results qualitatively match the old. But generally, a more accurate arithmetic model is desired.
- **Utilize GPUs and Parallelism:** The user's hardware includes high-end NVIDIA GPUs (RTX 3090 and 4090). The modernized simulation should exploit these to achieve better performance and scalability. Potential strategies include:
 - Using **vectorized numpy operations** to handle computations for all agents in bulk, and then employing **CuPy** (a GPU-accelerated NumPy-compatible library) to run those array operations on the GPU with minimal code changes. According to the CuPy project, in most cases it can be used as a *drop-in replacement* for NumPy, simply by importing `cupy` instead of `numpy` ¹⁴. This means we can prototype with NumPy (for correctness), then switch to CuPy to run on the GPU and gain significant speedups (often 10x-100x on heavy array computations ¹⁵).
 - Alternatively, use **Numba** (just-in-time compiler for Python) to accelerate loops with CPU vectorization or even target CUDA GPUs for custom kernels. Numba can translate annotated Python functions into optimized machine code. For instance, critical inner loops (like iterating over friends and goods to find exchange deals) might be good candidates for Numba, enabling parallel execution on multiple cores or offloading to the GPU if vectorization is complex.
 - We might also consider **PyTorch** or **JAX** to represent agent state as tensors and let these frameworks utilize GPUs. This is more complex (they are geared towards machine learning, but can be used for general compute on GPU). Given our needs, CuPy is likely the simplest path for GPU acceleration, since it aligns with the NumPy-based approach and doesn't force a paradigm shift.
- If extremely large-scale simulations become a goal (tens of thousands or millions of agents), we might explore specialized frameworks like **FLAME GPU 2**, a library specifically designed for large agent-based models on GPUs. FLAME GPU allows defining agents and their behaviors in a way that the system executes them in parallel on the GPU, achieving massive throughput ¹⁶. (It even has Python bindings where agent functions can be written in a restricted Python subset and transpiled to CUDA code ¹⁷.) However, using FLAME GPU would be a significant departure – essentially adopting a new engine – so our primary plan is to implement the logic ourselves in Python and only consider FLAME GPU if needed for extreme performance scaling.
- **Interactive Dashboard UI:** Instead of relying on static file outputs and offline analysis, we will build a live, visual interface for the simulation. The UI should provide:
 - **Real-time Metrics:** Plots of aggregate metrics over time (e.g. total production, average utility price, % of needs met, amount of spoiled goods, etc.), updated as the simulation runs. This gives an overview of the economy's state (e.g. detecting growth phases and cycles, as reported in the original results ¹²).

- **Agent Drill-down:** The ability to inspect individual agents or subsets of agents. For example, the user could click on or select an agent to see its current inventory, production output, trading partners, and maybe a history of its wealth or utility satisfaction over time. This addresses the original need to “study behavior and decision making of agents” ⁹, but in a more convenient way than reading debug text.
- **Regulatory Controls:** Tools for the user to act as a “regulator” in the simulated economy. For instance, UI sliders or inputs to impose a tax or subsidy on certain goods, adjust global parameters (like price elasticity or production efficiency factors), or toggle rules (e.g. enabling a market mechanism vs pure barter). The simulation should be designed to accept such interventions on the fly (at least between cycles) and reflect changes in agent behavior accordingly. This will allow experiments with policy: e.g. “What if we tax utility X heavily? How does the system respond?”
- **Strategic Agent Input:** Eventually, the user should be able to instruct one or more agents to follow specific strategies (“play a game”). In practice, this could mean overriding an agent’s decision-making AI with user-defined behavior. For example, an agent could be told to hoard a certain good, or to always reject trades below a threshold – and then one can observe how other agents respond and how the overall equilibrium shifts. Implementing this means our agent logic needs to be modular and overridable – possibly by plugging in a different decision routine for selected agents (the UI can provide a menu of strategy presets or even a scripting interface). This is a more advanced feature that will likely be built in a later phase, but we should design the software architecture with this flexibility in mind.
- **Modularity and Extensibility:** The new design should be modular. Each part of the simulation (e.g. agent state representation, consumption/production logic, trading mechanism, pricing, networking, etc.) will be a separate component or module. This encapsulation makes the system easier to understand and allows component-wise rewrites or upgrades in the future. For example, if we later want to replace the barter exchange algorithm with a different market mechanism, we should be able to swap out that module without rewriting the whole system. Similarly, the visualization layer should be decoupled from the simulation kernel so that the core logic can run headless (for batch experiments or automated testing) and only if a UI is attached do we engage the dashboard elements.
- **Maintain Output Parity and Scientific Validity:** It is important that the modernized simulation be validated against the original. The emergent behaviors reported (growth to a high level, economic cycles, distribution of “money” value vs. actual goods, etc. ¹⁸) should still manifest. We will verify that key metrics (like average utility price, total production, etc.) over time are consistent with the original model’s results (within reasonable variance given any stochastic elements). This entails careful testing and possibly reusing the original output logs as a baseline for comparison. In essence, while we might improve the internal implementation, the *economics* of the model must remain correct.

These goals introduce some challenges: ensuring Python can handle the workload, designing a responsive UI that doesn’t slow down the simulation excessively, and managing the complexity of rewriting a non-trivial simulation using AI assistance. Each of these is addressed in the plan that follows.

Recommended Technology Stack

Based on the above goals, we propose the following tech stack and tools for the modernization:

- **Core Language and Runtime: Python 3.x** (latest version) for the main simulation code. Python offers ease of development and a rich ecosystem, which is ideal for iterative, AI-assisted coding. Its dynamic nature will help during rapid prototyping with Codex (the AI can generate Python quite readily). We will use Python's standard scientific stack (NumPy, etc.) to handle calculations efficiently.
- **AI Coding Assistance: OpenAI Codex/GPT (via GitHub Copilot or ChatGPT)** for code generation. This will be used at multiple stages:
 - To translate portions of the C code logic into Python: We can feed Codex well-defined functions or code snippets from the C code with comments, asking it to produce Python equivalents. For example, we might prompt it with: "Here is a description of the `produceneed` function and what it should do... write it in Python." The assistant can often produce a good starting implementation.
 - To generate boilerplate code for class definitions, data structures, or repetitive sections: e.g. creating a Python class for agents with all needed attributes, or generating a loop that goes through all agents to perform an update.
 - *Important:* We will not rely on AI blindly. Each AI-generated piece will be reviewed and tested. We plan a **modular, incremental approach** (detailed later) so that we can verify correctness of each component as we build. This mitigates the risk of subtle bugs introduced by AI misinterpreting the requirements. Additionally, by using source control, we can track changes the AI suggests and revert if something goes wrong.
- **Data Structures and Agent Framework:** We will represent agents using either **Python classes** or structured NumPy arrays. Two possible approaches:
 - **Object-Oriented:** Define an `Agent` class with properties (needs, surplus, etc. possibly as NumPy arrays or lists of length 31 for each utility). Also define a `Market` or `Economy` class for global state (aggregates). Methods on these classes implement behaviors (e.g. `agent.produce()`, `agent.trade_with(friend)`, or static methods that handle multi-agent interactions). This approach is intuitive and mirrors the conceptual model (each agent is an object with its own state and behaviors).
 - **Vectorized State Matrices:** Use NumPy arrays to represent state of all agents collectively – for instance, a 2D array `needs[agent_index, utility_index]` for all agents. This is more akin to how the C code stored things (contiguously in arrays). It can make certain computations (like summing total production of a utility across all agents) very straightforward and fast (just a vectorized sum). It also is friendly for GPU acceleration, since one can move a whole array to GPU memory and operate on it in one go. The downside is that writing the logic in a vectorized style can be less clear (we have to think in terms of whole arrays rather than per-agent procedural logic), and some agent behaviors (especially discrete decisions or conditional logic) are harder to vectorize.

A hybrid is likely: use classes to organize code, but internally use NumPy arrays for heavy data fields. For example, `Market` could hold big arrays for all agents' inventories, while each `Agent` object holds indices or slices into those arrays. However, managing consistency could get tricky. Alternatively, we might start with the simpler OOP approach for clarity, then profile and gradually refactor bottlenecks

into vectorized operations. **Mesa** (discussed next) uses an OOP approach (each agent is an object with a `step()` method), which is natural but may not leverage the GPU easily. We will choose a structure that balances clarity with performance.

- **Agent-Based Modeling Framework (Optional):** Consider using **Mesa**, a Python ABM framework ¹⁹. Mesa provides a well-defined structure for agent-based simulations, including a scheduler for agent activation and built-in tools for data collection and visualization. It features a browser-based visualization server that can produce interactive dashboards for the simulation (with charts and even live agent views on a grid or network) ²⁰. Key reasons to use Mesa:
 - It enforces modular design (separating Model, Agent, and Schedule).
 - It has an existing **interactive dashboard** system: one can create a **Visualization** with components like charts or a network graph of agents, and Mesa will handle updating them as the model runs. For example, Mesa's tutorial demonstrates building a simple money exchange model with a live dashboard ²¹ ²⁰.
 - Mesa is pure Python (no GPU support by itself), but one can potentially integrate NumPy/CuPy calculations within agent methods. There's also an extension called **Mesa-Frames** for better performance and possibly parallel execution ²², and a **batch runner** for experiments. Mesa could save time in creating the UI and structuring the project.

However, Mesa might impose some limitations: it might not scale to thousands of agents easily unless optimized, and using GPU with Mesa's step-by-step loop may require custom integration. If Mesa's benefits outweigh its constraints, we should use it for rapid development of the UI and model structure. Otherwise, we can implement our own lightweight loop and perhaps use **Plotly Dash** or similar for visualization.

- **Visualization and UI Framework:** For a dashboard-style UI in Python, two major options stand out:
 - **Plotly Dash:** Dash is a popular open-source framework for building interactive web dashboards entirely in Python. It's built atop Flask (for the web server), React.js (for front-end components), and Plotly.js (for charts) ²³. With Dash, we can create an application that displays live graphs of our simulation and has interactive controls (sliders, dropdowns, buttons) that the user can adjust. The simulation can run in the background (possibly as a thread or via periodic callback updates). Dash apps run in a web browser, which is ideal for a modern UI. The advantage is full flexibility – we can design any layout, use rich chart types, and even render interactive network graphs (Dash has a Cytoscape component for network visualization of agents). Dash would also facilitate eventual deployment as a web app (it can be hosted on a server or even shared via services like Heroku or Dash Enterprise).
 - **Mesa's Browser GUI:** If we use Mesa, it provides a simple web server that can show the simulation. We can define visualization elements: e.g., a `NetworkModule` to display the social network graph of agents, `ChartModule` to plot variables over time, and `Text/HTML` for agent details. The Mesa GUI isn't as customizable as Dash, but it covers basic needs and is tightly integrated with the model (it can call our model's data collectors each step). Mesa's *Visualization Tutorial* provides guidance for making such dashboards within the Mesa framework ²⁴ ²⁵. Using it could be faster to get something running, though perhaps less flexible than Dash for very custom controls.

Given the desire for a rich and possibly unique UI (with regulator controls, etc.), **Dash (Plotly)** is likely the top choice if we don't go with Mesa. It gives us full control over the interface. We will need to manage the link between the simulation and the UI manually (e.g., using Dash's callback system to step the simulation and retrieve data for graphs). A possible architecture is to run the simulation in intervals (e.g. simulate one cycle per second, or as fast as possible but yielding to update the UI periodically), and have Dash callbacks fetch the latest metrics to update charts. Alternatively, run the simulation in a

background thread continuously and use WebSocket or interval polling to update the UI. We must ensure thread-safety if doing this (since Dash is event-driven). Another library to mention is **Streamlit**, which easily creates dashboards in Python, but it's more suited to static analysis or slower update cycles (it reruns script on each interaction). Dash, on the other hand, is designed for dynamic updates and could be made to show near-real-time changes.

In summary, we recommend **Dash for a custom interactive dashboard**, unless Mesa is adopted for the core (in which case use Mesa's built-in viz to start, potentially moving to Dash if needed for advanced control).

- **GPU Libraries:** As discussed, **CuPy** is a prime candidate to leverage the RTX 3090/4090 GPUs. We will use CuPy for heavy numeric arrays if profiling shows CPU NumPy is a bottleneck. Installing CuPy is straightforward (pip packages available for CUDA 11.x and 12.x ²⁶). By maintaining compatibility between our code and NumPy's API, switching to CuPy is mostly a matter of adding `import cupy as np` (or a slight refactor where we can easily toggle which library is being used). We should also ensure that intermediate data transfer between CPU and GPU is minimized (ideally, once data is on the GPU, keep it there for the duration of simulation steps).

Additionally, for tasks not easily handled by built-in array operations (like if we need a custom kernel to iterate over agents in parallel), CuPy allows writing custom CUDA kernels in C++ and integrating them ²⁷. We might avoid that complexity initially, but it's good to know it's available for fine-tuning performance-critical sections.

- **Parallel Processing:** Beyond GPU, we should exploit CPU parallelism if possible. Python's Global Interpreter Lock (GIL) makes true multi-threading difficult for CPU-bound tasks, but we have options:
- Use **multiprocessing** (multiple Python processes) to divide the simulation if needed (e.g., simulate different segments of agents in separate processes or run multiple scenarios in parallel). This requires careful consideration of how to partition the model, since agents do interact (so splitting one simulation across processes is non-trivial unless using shared memory for global state).
- Use **Numba's** `prange` for automatic parallelization in JIT-compiled loops. If we JIT compile an update loop with Numba and specify `parallel=True`, it can vectorize and multi-thread the loop internally (release the GIL and use multiple cores).
- Ensure any heavy library routines (NumPy, CuPy, etc.) are using native parallelism. NumPy linear algebra often uses BLAS under the hood which can be multi-threaded. CuPy (CUDA) will naturally use thousands of threads on the GPU hardware. Thus, by delegating work to these libraries, we implicitly get parallel execution outside of Python's single thread.
- **Agent Networking & Data Structures:** We will likely use **NetworkX** (Python library for graphs) or simple adjacency lists to represent the social network (friend connections between agents). This can help if we want to visualize or analyze the network structure (e.g. finding clusters). However, for simulation logic, a lightweight structure (each agent having a list of friend indices) is sufficient, as in the C code. Maintaining and updating this network (agents randomly adding new friends, etc.) will be part of the simulation loop. We should keep it efficient; e.g., if using Python lists or sets for friends, ensure operations are not too slow. Given the network size is at most 10 friends per agent in the original (MAXGROUP=10), this is not large and not a bottleneck.
- **Visualization of Network:** If we want to show the social network in the UI, Dash's Cytoscape component can be used to draw a graph of agents (nodes) and friendship links (edges). We

should allow highlighting of one agent and its immediate network, for drill-down analysis. Alternatively, a simpler approach is to use a heatmap or matrix to show interactions or to just list an agent's friends and their trade frequency.

- **UI Design Tools:** For plotting metrics, we will use **Plotly** (which underpins Dash graphs). For example, a time series of “average utility price” can be a Plotly line chart that updates each simulation step by appending a new point. We might also include bar charts or histograms for distribution of something (like distribution of wealth among agents at a given time). For agent details, Dash can display tables or text fields. The UI should be arranged in a dashboard format – perhaps with a top section for global metrics, a side panel to select or input interventions, and a bottom or modal section for detailed agent info.
- **Source Control: Git** with a remote repository on **GitHub** (or GitLab/Bitbucket as needed). This is essential for tracking the rewrite progress and collaborating with AI effectively. We will:
 - Keep the legacy C code and design notes in the repository for reference.
 - Create a new Python project structure (perhaps using a cookie-cutter or standard layout: e.g. a main package for simulation code, a separate module for the Dash app, etc.).
 - Commit frequently, especially after each module is generated by AI and passes basic tests. Commits serve as checkpoints – if AI suggestions go awry, we can rollback to a stable state.
 - Use branches for major steps: e.g., a branch for initial translation of core logic, another for integrating the UI, etc., merging into a `main` or `dev` branch once each part is stable.
 - Possibly set up **GitHub Actions** for continuous integration testing – for instance, run a quick simulation and verify outputs to catch regressions.
 - If multiple developers (or instances of AI assistance) are working, issues and pull requests can manage changes. Even as a single developer project, writing clear commit messages (possibly AI-assisted via tools like GitHub Copilot's CLI) will document the evolution.
 - GitHub integration also makes it easier to later showcase the project or open-source it, if desired, and to integrate with project management tools (Kanban boards for tracking tasks, etc.).

In summary, the technology stack centers on **Python** for core logic, **AI assistance** for development, **NumPy/CuPy** for performance, and **Dash or Mesa** for visualization. This combination should meet the goals of improved maintainability, high performance (utilizing GPUs), and rich interactivity.

Proposed Architecture and Module Design

Before coding, it's crucial to outline the new system's architecture. A clear separation of concerns will guide the AI-generated code and ensure we can extend the simulation with new features (like user interventions) easily. The proposed architecture includes the following main components:

1. **Simulation Core (Engine):** This is the heart of the model – responsible for iterating through time *steps* (cycles). It will likely be implemented as a loop that updates all agents and the market for each period. We can encapsulate this in a `Simulation` or `Economy` class. This class would contain:
 2. The collection of all agent objects (or arrays of agent state).
 3. Global parameters (some might be adjustable by the user through the UI, e.g. a flag to enable taxation).
 4. The main `run_step()` or `run_period()` method that executes the sequence of updates that correspond to one cycle (period). From the original logic, a single period might include:
 1. Each agent consumes from surplus to meet needs (reducing surplus, satisfying some needs) ²⁸ .

2. Agents engage in **barter exchanges**: finding mutually beneficial trades with friends to cover remaining needs. This involves the complex negotiation algorithm (finding the best exchange deal for a need among friends) ²⁹ and executing trades (transferring surplus goods between agents) ³⁰. The code has functions like `getfriendexchangeindexes` and `executeexchange` for this.
3. Production phase: agents produce goods to meet their unmet needs. There might be multiple rounds – one for primary needs, possibly another for small needs, and a surplus production round if they still have time capacity ³¹ ³². Also a “leisure” production if they have free time (producing any extra if capacity remains) ³³.
4. Post-production adjustments: update each agent’s experience – adjust prices they’re willing to buy/sell at based on trades (functions `adjustpurchaseprice` and `adjustsalesprice` apply heuristics) ³⁴, update their social trust/transparency levels with friends (function `calibratefttransparency`) ³⁵, and possibly form new friendships (the C code had logic to randomly add a new friend connection occasionally).
5. End of period bookkeeping: agents finalize their stocks (calculate what spoiled due to excess beyond storage limit, etc.), and global metrics are computed (totals, averages). The function `endmyperiod` in C handled updating production cost averages, efficiencies, and marking if an agent ran out of time (periodfailure) ³⁶.
6. Advance to next period: global time period counter increases, and some history memory (like what was “recently produced” vs older history) is rolled over for the next cycle.

The new simulation core will implement these steps, likely as separate methods for clarity (e.g. `do_consumption()`, `do_exchange()`, `do_production()`, etc., all called within one period update). This corresponds closely to the original structure (which had separate functions for each stage, invoked in sequence in `evolution()` ⁵). We will preserve that high-level order, but with cleaner abstraction boundaries. For example, within the exchange step, we can break down the sub-tasks (evaluate possible trades, pick best trade for each agent, execute trades) so it’s easier to modify or optimize them later.

1. **Agent Model:** Each agent can be an object or a structured record. We’ll define the agent’s state properties clearly:
2. Permanent attributes: e.g. *efficiencies* for each good (how effectively they produce it), *giftedness* flags (some agents start with higher efficiency in certain goods), and possibly an ID and maybe a category (if we later classify agents).
3. Dynamic state per cycle: *current need* for each good, *current surplus* of each good, *stock limits* (how much they can store), *current stock* (though stock might be same as surplus here, since they produce and hold until trade/consume), *time remaining* in the period (to constrain how much they can produce).
4. Social network info: list of friend agent IDs, and for each friend some relationship data like *transactions count* and *transparency/trust* per good (the original `EVENTS` struct stored how many units agent bought/sold with that friend and a “transparency” metric ³⁷).
5. Pricing expectations: arrays for *sales price* and *purchase price* thresholds for each good (these are the agent’s private valuation thresholds, updated as they learn).
6. Statistics accumulation: counters for how many times they bought or sold each good this period, sums of values traded (used to adjust prices), etc.

We will encapsulate these in the Agent class. Some of these might remain NumPy arrays of length 31 for efficiency (e.g. storing needs or prices), but accessed via properties.

Additionally, agent methods will implement behaviors: - `consume_surplus()`: have the agent use their surplus to satisfy their own needs up to possible, reducing need and surplus. -

`evaluate_exchange_options()` or similar: have the agent figure out what it wants to trade. In the C code, for each unmet need, the agent looked at each friend's surplus to see if a trade can be made (I give X, get Y). This might produce an index or score for each possible exchange. The logic was fairly complex, ensuring the chosen trade is most valuable and meets both parties' criteria. We may simplify or at least clearly separate the concerns (perhaps first find all friends who have the item I need and see what item they might want that I have surplus of, then evaluate some utility value). - `produce_goods()`: logic for producing goods given remaining needs and time. Possibly split into primary production (produce what is needed) and secondary (use leftover time to produce extra of gifted skill or for trading). - `update_prices()` and `update_trust()`: adjust the agent's internal thresholds and relationship transparencies based on what happened this cycle (e.g. if a friend consistently trades fairly, trust increases; if an agent ended with unmet needs or oversupply, they adjust willingness to pay/accept prices). - If implementing agent strategies, we might have a method `decide_actions()` that could be overridden by subclasses or modified by UI input. For example, a subclass of Agent could be `GreedyAgent` or `HumanControlledAgent` with different trade behavior. The simulation core would call a generic interface (like `agent.step()` which internally calls the series of consume, trade, produce, etc.), and a specialized agent could override parts of this process.

1. **Global Market Model:** Even though this is a decentralized barter simulation, the original code had a `market` struct to track global aggregates and enforce some constraints (like price elasticity adjustments). We will have a `Market` or `EconomyStats` component that computes and stores:
 2. Total production of each good in the period, total consumption, total spoils, etc.
 3. Average price of each good (they computed this as average production cost weighted by output, to derive an emergent "money" metric) ³⁸ ³⁹ .
 4. Number of agents acting as consumers/producer/retailers for each good (the roles count, tracked each cycle) ⁴⁰ .
 5. Possibly the emergent "value of money" if any (in the original, one good might become a de facto medium of exchange – they observed that the "most common exchange medium" had a value that stabilized, indicating emergent money).
 6. Period counter, etc.

The simulation core will update these global stats at the end of each cycle (similar to the C function `evaluatemarketprices()` which recalculated average prices and other stats ⁴¹). These stats are what we'll feed to the UI charts as well. By separating this out, if we later introduce a centralized intervention (like a government agent or external shock), we can incorporate it here or as an external module.

1. **User Interface Module:** The UI (if using Dash) will be a separate part of the code, possibly in its own script or section. Its job is to launch a Dash server with layout and callbacks:
2. **Layout:** Define HTML/React components like graphs (`dcc.Graph` for time series plots), dropdowns for selecting an agent to inspect, sliders for adjusting tax rates or other parameters, and perhaps text boxes or tables for showing an agent's inventory. Layout can be divided into sections: e.g. *Global Metrics Panel*, *Agent Detail Panel*, *Controls Panel*.
3. **Callbacks:** Functions that respond to user input or timer events. For example:
 - A callback that triggers every *n* milliseconds (using `dcc.Interval`) to advance the simulation by one step and then update the graphs with new data. Alternatively, a manual "Play/Pause" control to start or stop the simulation loop.
 - A callback for when the user selects a different agent (update the agent detail panel with that agent's latest data).
 - Callbacks for control changes: e.g., if the user moves a slider to set a 5% tax on Utility 3, that callback will update a parameter in the simulation core (like

`market.tax_rates[3]=0.05`), which the simulation uses in subsequent steps to tax trades or reduce production of that good accordingly.

- A button to execute a one-time intervention (e.g. “Inject surplus good 5 to agent 10” as a test scenario, or “Reset economy”).

The UI module will need access to the simulation state. We can handle this by running the simulation in the same process (the Dash callbacks can directly call methods on a `Simulation` object that’s kept in memory). One complexity is making sure the simulation loop doesn’t block Dash’s event loop. We might run the simulation in a background thread and have it periodically emit updates (using a thread-safe queue or so). Or simpler, use the Dash interval callback as the clock: each tick, call `simulation.run_period()` once. This essentially freezes the UI update rate to the interval (which might be fine if we choose, say, 2 updates per second). If the simulation is fast, we could even batch multiple cycles per interval. We’ll decide the approach based on performance testing.

If using Mesa’s server instead: Mesa runs the model step by step with a built-in loop and uses a simple synchronous update to the browser via websockets after each step. We would follow Mesa’s paradigm (the code would define a `step` method for the model and agent, and Mesa handles UI refreshing). Mesa’s UI components (charts, etc.) would be defined in code and they automatically call the model’s data collectors. This is somewhat more restrictive but less code for us to write. We will evaluate if Mesa can easily incorporate our planned controls (Mesa does allow adding custom buttons and elements, but it might require some JavaScript coding for complex interactions). If we need something highly custom, Dash is preferable.

1. **Configuration and Persistence:** Provide a way to configure simulation parameters (population size, initial conditions, etc.) easily. This could be a JSON or YAML config file, or just a Python dictionary. Also, integrating with source control, we can store different scenario configurations. Additionally, if we want to save simulation state or results, we should have a mechanism (e.g. the ability to dump all agent data to a file or database at the end of a run, or periodically). This isn’t a primary requirement, but for scientific completeness it’s good to allow reproducibility. For instance, logging every period’s summary to a CSV (similar to what the original did in `adamsmith.txt`) can be an optional feature to enable in headless runs.
2. **Extension Hooks:** Design the architecture with extension in mind. For example:
 3. A plugin system for new “policies” or events. Perhaps define an interface for regulatory interventions that can be applied each period (like functions that modify agent state or transactions if active). Then the UI controls simply toggle these on/off or adjust parameters. This way adding a new policy (say, universal basic income distribution each period, or a shock that halves all agents’ surplus every 50 periods) can be done by writing a new function and hooking it into the simulation loop without altering core logic.
 4. Strategy override for agents: we might implement this by giving each agent a strategy object or a set of flags. The agent’s decision-making methods consult those: e.g., if `agent.strategy == "Greedy"`, maybe skip certain trades or always produce the most expensive good, etc. For a user-controlled agent, we could allow the UI to directly set that agent’s next action (like a manual override: produce X of good Y, regardless of its algorithm). We need to ensure the architecture allows an external command to influence an agent’s state safely (perhaps via the simulation core issuing a command to that agent before the cycle step begins).

In summary, the architecture centers on a **Simulation Engine** managing a collection of **Agents** and a **Market** state, with a clean API for the **UI** to interact (start/stop simulation, inspect/modify state). The design emphasizes modular separation: one could run the simulation without the UI (for automated

tests or batch runs) by simply not instantiating the dashboard. This separation will also help in testing (we can write unit tests for agent behaviors and for one cycle of simulation logic in pure Python without needing the UI).

Incremental Development Plan (Using AI Assistance)

To manage the complexity of the rewrite, we will proceed in stages, rewriting and testing one component at a time. This also aligns with effectively using Codex/AI, as we can focus the AI on one well-defined task at each step. Below is a phased plan:

Phase 0: Preparation and Understanding

- **Read and Annotate Legacy Code:** Thoroughly go through *EmergentMoney.pdf* and comment important parts of *adamsmithfeb17B.c*. We will document what each section does in plain language. This is partly done in this report, but we might create a more detailed mapping (e.g., "Function X in C corresponds to Y in our planned design"). This serves as a reference for prompting Codex and for verification of correctness. The document already gives clues like "*function X calculates exchange options for a needed good*". We will translate such insights into comments or pseudocode.

- **Set Up Tools:** Ensure we have access to Codex (via the OpenAI API or GitHub Copilot in an IDE) and that our development environment can run and test the code (install Python, needed libraries like numpy, etc., and possibly set up a repository in GitHub so Copilot has context if needed).
- **Plan Data Structures:** Decide on core data representation (object vs arrays). For start, we might implement in a straightforward way (Agent objects with Python lists for needs/surplus) and later optimize. Document this choice so that while prompting AI we can say "we have an Agent class with these attributes..."

Phase 1: Define Data Models

- **Agent Class and Market Class:** Using AI assistance, write the skeleton of the `Agent` class and `Market` class in Python. This includes defining all relevant fields (perhaps use Python `dataclass` for brevity). For example, we'll prompt Codex: "Create a Python dataclass for Agent with fields: *id* (int), *needs* (list of floats length *N_UTILS*), *surplus* (list of floats length *N_UTILS*), *efficiency* (list of floats length *N_UTILS*), *purchase_price* (list of floats), *sales_price* (list of floats), *friends* (list of friend IDs), etc.". The AI should generate a reasonable class definition. We will then manually adjust any naming or types as needed. We ensure that default initial values make sense (e.g. *needs*=0, *surplus*=0 for all goods initially except maybe some initial surplus as in the C init).

- **Global Parameters Enumeration:** Create a config structure or at least module-level constants for things like number of agents, number of goods, initial values (similar to those `#define` in C). This ensures the code isn't littered with magic numbers. For example, `N_AGENTS = 300` (initially), `N_UTILS = 31`, initial efficiency values, etc. These can be stored in the Market class or a separate Config class.
- **Initialization Functions:** Implement `initialize_market()` and `initialize_population()` akin to the C code's `initmarket` and `initpopulation`⁴². We can use Codex to translate the logic: it should set each agent's initial needs, surplus, etc., and compute the market's initial period length (total needs) and leisure time, etc. We'll verify this matches the original logic (for example, ensuring that initial surplus equals basic need production as a starting condition⁴⁴). This also involves assigning random gifted skills to

agents (the C code's `selectgiftedpopulation` randomly flags some talents). We might use Python's random module to do similarly. We will set a random seed for reproducibility.

- **Basic Data Integrity Test:** After initialization, we will run a quick check (write a small test or just inspect) to confirm things like: total initial needs equals market.periodlength, each agent's surplus equals their basic need, etc., matching the C code's expectations. This ensures our data model and init are correct.

Phase 2: Implement Core Simulation Steps

We proceed to implement the main cycle actions one by one, testing as we go:

- **Consumption Step:** Code the agent method (or standalone function) for `consume_surplus()`. This should iterate through each good and if the agent has need > 0 and surplus > 0 of that good, transfer as much as possible from surplus to fulfill need. We can use AI to generate this loop easily. After consumption, the agent's need for that good should reduce (possibly to zero) and surplus correspondingly reduces. Also track if any need remains. We test this on a simple scenario (e.g., agent with need 5 and surplus 3 of a good ends with need 2, surplus 0). This step is straightforward and mostly bookkeeping.
- **Trade (Exchange) Step:** This is the most complex part. We might break it down:
 - *Friend exchange evaluation:* For each agent, for each remaining unmet need, evaluate potential exchanges with each friend. The C code did something like double-nested loops (need, friend, possible gift goods) to find the best trade ⁴⁵ ⁸. We can simplify by structuring it as: each friend can offer some goods (their surplus) and will want something in return (something they need that our agent has surplus of). We might attempt a heuristic approach: for a given agent A and friend B, find one pair of goods (X from B to A, Y from A to B) that is mutually beneficial and gives the best value to A's need. We can translate the logic from the original function `getfriendexchangeindexes` ²⁹. Possibly use AI to get a starting point but be ready to refine it.
 - We might not implement the entire exchange in one go. Instead, aim to get a basic barter mechanism working (even if simpler than original) and then refine. For instance, initially implement: each agent looks at each friend and if the friend has surplus of something agent needs and agent has surplus of something friend needs, do an exchange of one unit each. This is simplistic but we can use it to test that trades occur. Then we can elaborate to account for varying values and larger quantities as per original logic.
 - There is also a question of *synchronization*: do all agents trade simultaneously, or one by one? The original code likely had each agent in turn attempt trades (which could affect others' surplus by the time it's their turn). This order dependency can matter. We might preserve the sequential approach for simplicity: loop over agents, and for each, attempt to fulfill one need via trade (the best trade available). This could introduce bias (earlier agents get first pick), but that might have been the case originally as well. Alternatively, we could collect all potential trades and then execute them in some order. However, to keep it simple and similar to the original, sequential is fine.
 - **AI usage:** We can prompt Codex with something like: *"Implement a function where agent A finds the best trade with its friends. Each friend B has some surplus goods and some needs. A trade consists of A giving good g1 and receiving good g2 from B. The trade is beneficial if g2 fulfills A's need and g1 fulfills B's need. We choose the trade that maximizes A's utility index = (amount of need fulfilled) * (some value factor)... etc."* This is complex, so we might feed parts of the original algorithm as guidance. We should be careful to verify the AI's logic matches our intentions. Likely we will need

to adjust and test thoroughly (maybe create a scenario with two agents and ensure the trade logic does what we expect).

- **Trade Execution:** Once a trade deal is decided (e.g. A will give X units of good i to B, and B will give X units of good j to A), implement a function to execute it: subtract from A's surplus and B's need, etc., and update records (like increment transactions count, etc.). Mind that goods might have different "values"; the original allowed unequal exchanges if values differ, but they often traded equal time cost units. For now, we can trade unit-for-unit or proportional to some price. We will later refine using the price thresholds (if implemented).
- **Validation:** After implementing trading, run a small test: e.g., two agents with complementary needs and surpluses and see if they swap. This will confirm that at least simple barter works. We should also verify that no conservation laws are broken (goods shouldn't appear or vanish incorrectly – except spoilage or production by design).
- **Production Step:** Implement functions for agent production. The original had a concept of each agent having a time budget per cycle (periodlength divided among all goods needs) and producing goods to meet their needs first, then possibly producing extra (surplus) if gifted or if they have leftover time (leisure production) ³³.
- Start with a simple model: each agent, for each unmet need, produces that good until the need is met or until they run out of time. Efficiency comes into play: an agent with efficiency e in a good means producing one unit consumes 1/e of their time unit. So if e=1.0, one unit costs 1 time; if e=2.0 (gifted in that skill), one unit costs 0.5 time, etc. We should implement time consumption and not allow an agent to exceed their `period_remaining` time.
- After meeting needs, if they have time left (free time), they could produce more of whatever they are best at (surplus to trade). The original simulation allowed "surplus production" focusing on goods that have trading potential – likely goods where they had high efficiency and others need them. We can approximate that by: agent uses remaining time to produce extra of goods where they are producer/retailer (role) or where price is high. Initially, maybe just produce more of their most efficient good as surplus.
- We should update agent's surplus inventory and zero out their needs if fulfilled. Also update their `period_remaining` time as they produce.
- Use Codex to speed up writing these loops with the given constraints, then test on an agent scenario (like one agent with needs and given high efficiency in one good – it should produce that good quickly to fulfill need).
- Also, update global production totals and maybe track cost if needed (like sum of squared inverse of efficiency for cost, though that might be for price calculation).
- **Price and Value Update:** Implement functions akin to `adjust_purchase_price` and `adjust_sales_price` from the C code ⁴⁶ ⁴⁷. The idea is: after each period, an agent looks at what happened – did they manage to buy what they needed? Did they fail to sell surplus? Based on that, they adjust their personal value estimate for goods. For example, if an agent had to go without some utility, they might raise the price they're willing to pay for it next time (indicating it's more valuable to them). Conversely, if they have surplus they couldn't get rid of, they might lower the price they demand for that good. These heuristics mimic learning and drive the system's price convergence.
- We can derive heuristics from the C code comments and structure. The code comments mention things like: *"if one becomes wealthy, needslevel increases... if less wealthy, needslevel reduces"* ⁴⁸ and *"adjust purchase prices – heuristics for purchaseprice"* ³⁴. We will likely implement simplified

rules: e.g., for each good, if agent ended the period still needing that good (need > 0 unsatisfied), increase `purchase_price[good]` by a small factor (up to some max). If agent ended with surplus of a good unsold/unused, decrease `sales_price[good]` slightly (so they'll be more willing to trade it away next time).

- Transparency/trust update (`calibrateftransparency` in C) adjusts how much an agent trusts each friend for each good, likely based on whether trades happened or if friend failed to deliver promised trades. We might for now assume a simple model: if a friend traded successfully, increase trust; if a trade failed or friend had nothing to offer, maybe decrease a bit. This can be elaborated later; the impact of this is to influence who they prefer to trade with (more transparent friends are favored).
- **Cycle Integration:** Once the above pieces are in place, implement the `Simulation.run_period()` that calls them in sequence for all agents and then updates global stats. This will look like:

```
def run_period(self):
    for agent in self.agents:
        agent.consume_surplus()
    for agent in self.agents:
        agent.find_and_execute_trade() # (which internally loops over
needs and friends)
    for agent in self.agents:
        agent.produce_goods()
    for agent in self.agents:
        agent.update_prices_and_trust()
    self.market.compute_global_stats(self.agents)
    self.market.period += 1
```

This is a conceptual outline. In practice, we might combine or tweak orders (e.g., the original did consumption, then potentially multiple rounds of trade and production – basic need production, then surplus trading, then surplus production, then leisure production). We can start with one round of each and see if that yields plausible results, then consider adding a second “surplus trade & production” round if needed (to mimic the original’s iterative convergence within a cycle).

Use the AI to generate parts of this high-level loop after we provide it context of what each step does (perhaps via docstrings on those methods).

- **Testing Iteratively:** At this stage, we can try running a few cycles with a very small number of agents (maybe 5 agents, 3 goods) to see if everything runs without errors and the data changes make sense (e.g., needs go down to zero by cycle’s end, surplus distributed logically, etc.). We likely will need to debug some logic (the trade algorithm especially). We can add printouts or use a debugger to trace a single cycle.
- **Benchmark vs Original (small scale):** If possible, run the original C code for a couple of cycles with a tiny population (maybe in a controlled scenario) to gather reference output. If compiling and running is easy, we could feed it a fixed random seed and capture certain metrics (like total surplus, total spoils after 1 cycle) and compare to our Python version on the same seed. However, given the complexity, we might instead rely on qualitative validation: ensure, for example, that in our sim initially everyone meets their needs by own production, then as social

network forms, specialization and trade begin – which mirrors the report’s description of initial results ⁴⁹ .

Phase 3: Performance Tuning

- Once the basic simulation loop is working correctly for small cases, we profile it with a larger case (e.g., the default 300 agents, 31 goods, run 50 cycles) using Python’s profiling tools. Identify bottlenecks: likely the trade-finding loop is expensive ($O(N_{\text{agents}} * N_{\text{friends}} * N_{\text{goods}}^2)$ potentially). Also production loops over goods. We then address these: - See if we can vectorize parts: for instance, computing total production per good is a simple sum across agents – make sure we use NumPy for that rather than Python loops ⁵⁰ ³⁹ . Similarly, if adjusting all agents’ prices by some rule, we can do that with array ops. - If the trade logic is too slow in pure Python, consider using **Numba** to JIT it. We can mark the inner loops as `@njit` (NoPython mode) which will compile it to C-like speed. We have to be careful with data structures though – might use NumPy arrays for input/output to the Numba function. - Alternatively, try converting some parts to use matrix operations. For example, one could represent needs and surplus of all agents as matrices and compute a matrix of potential trades (this is complex because of conditions, but maybe possible for some simpler exchange model). - Use multi-threading if possible: If we stick to Python loops, maybe use the `concurrent.futures` to parallelize by agent (e.g., split agents into 4 groups and have 4 threads do trades simultaneously). Because agents interact (through trades), this is tricky – we would need to partition by something like friends or ensure threads don’t conflict. Possibly not worth the complexity at this time. - Use GPU: Once correctness is confirmed, swap in CuPy in place of NumPy for heavy computations. For example, if we maintain a NumPy array of needs of shape $(N_{\text{agents}}, N_{\text{goods}})$, moving that to CuPy and doing operations (sums, comparisons) could speed up parts of the simulation. The trade step might not easily map to GPU if it involves a lot of branching logic, but maybe we can offload sub-calculations (like evaluating all possible exchange values could be done in a vectorized manner on GPU if framed as matrix ops). This is an advanced optimization; we should ensure the simulation works with CPU first, then profile to see if GPU usage is beneficial (given $N=300$ or even 3000 , a single GPU might actually not be fully utilized unless we increase agent count by an order of magnitude; but if we do, the GPU will shine). - Memory considerations: The original mentions the data can be gigabytes if one stores everything ⁶ . We won’t store full history by default (just current state, and maybe aggregate stats history for plotting). $3000 \text{ agents} \times 31 \text{ goods}$ with a few arrays isn’t huge (a few hundred thousand numbers). Even $100k$ agents would be manageable memory-wise (tens of millions of numbers, meaning on the order of a few hundred MB). So memory should be fine on modern PCs, but we should avoid unnecessary copies. Using NumPy/CuPy helps, since operations happen in-place or in efficient C loops.

- Re-run the simulation after optimizations to see if we can approach real-time speeds (if interactive, ideally one cycle computes in significantly less than 1 second for a few thousand agents). If not, adjust expectations (maybe only simulate smaller numbers in real-time; for larger, run faster-than-real-time in headless mode). In any case, document the performance and ensure it’s sufficient for the intended use.

Phase 4: Develop Visualization and UI

- If using **Dash**: Begin building the dashboard once the simulation backend is solid. This can be done in parallel with performance tuning, since the UI doesn’t fundamentally change the simulation logic. - Create a Dash app layout with placeholders for graphs and controls. For example, a line chart for “Average Utility Price vs Time”, another for “Total Production and Spoilage vs Time”, etc. Also a dropdown to select an agent and display its inventory and maybe a small network graph of its friends. - Implement callbacks: - A callback tied to an Interval that calls `simulation.run_period()` and then updates the data for each graph. We have to store simulation outputs in a place accessible to callbacks – Dash provides a concept of “Store” (an invisible component holding JSON data) or we can use a global `simulation` object (since Dash callbacks can close over outer scope if not using multiprocessing mode).

We'll likely keep a global simulation instance that the callbacks use. - A callback for agent selection: when user picks an agent ID from dropdown, update the agent detail panel (which could show e.g. a bar chart of that agent's current surplus of each good, or text of their utility fulfillment percent, etc.). - Callbacks for interactive controls: e.g., a slider for tax on Good X: when moved, set `simulation.tax_rate[X]` and perhaps display the value. The simulation loop will read this and adjust trading or production behavior (for instance, if a tax on a good is set, maybe when computing trades, if that good is being given by someone, we reduce the utility by the tax rate to simulate losing some to tax). - Buttons: "Pause/Resume Simulation" (to stop or start the periodic updates), "Step One Cycle" (to advance manually), "Reset Simulation" (to reinitialize). These are useful for user control. - Use the AI to assist with writing the Dash callback boilerplate, as it can be finicky. Ensure to test the app manually – run the server and play with the controls to see if everything updates. We might find we need to throttle updates or handle the case when simulation is paused, etc.

- If using **Mesa**: Set up a Mesa Model class (inheriting `mesa.Model`) and Agent class (inherit `mesa.Agent`). Transfer our logic into Mesa's paradigm (Mesa calls `step()` on each agent; we might instead use a `RandomActivation` schedule to call agent methods in random order). Add `DataCollector` for collecting global stats each step. Then configure Mesa's visualization:
- Define chart modules: e.g., `ChartModule([{"Label": "AveragePrice", "Color": "Blue"}, ...], data_collector_name="datacollector")` to plot average price over time.
- Define a text element or custom module to display selected agent info (this might require custom JavaScript if we want dynamic selection, which is a bit advanced; Mesa's default doesn't have a dropdown for agent selection out-of-the-box).
- Mesa does have the concept of user-settable parameters at launch, but interactive runtime controls might be limited. Possibly we can add a custom button that triggers a function (Mesa's server allows some custom buttons that call a method on the model).
- If Mesa's capabilities suffice for a basic dashboard, we might use it initially (especially if we want results quickly). But given the specific asks (drill-down, regulator controls), we suspect a custom Dash app will offer a better experience.
- In either UI approach, keep an eye on performance. UIs can slow things down if updating too often or drawing too much. We may need to update charts only every few simulation steps, or limit how much history is plotted (e.g., last 100 cycles). For agent network visualization, 300 nodes is okay, 3000 might be heavy to draw; maybe only show a subgraph of interest.

Phase 5: Testing & Validation

Testing will be woven into earlier phases, but here we formalize it: - **Unit Tests:** Write tests for critical functions: - Consumption: ensure an agent with known needs and surplus results in correct post-consumption state. - Production: test that an agent with certain efficiency and available time produces the expected amount of goods. - Exchange: set up two agents with a simple scenario where a trade is obvious (A needs 5 of good1, B has surplus 5 of good1; B needs 5 of good2, A has surplus 5 of good2) and verify after trade both needs are satisfied and surplus decreased accordingly. - Price adjustment: simulate a scenario where an agent couldn't buy a needed item and see that their `purchase_price` for it increases. We can use Python's `unittest` or `pytest` frameworks. These tests help ensure that as we tweak performance or add features, we don't break basic behaviors.

- **Integration Tests:** Run the full simulation for a short time and validate aggregate outcomes:
- For example, start with all agents identical and no social network, run 1 cycle – we expect each agent just produces for themselves and no trade occurs, so everyone meets needs and no one accumulates surplus or spoilage.

- Introduce one gifted agent who is very efficient in one good, and check that over a few cycles, that agent starts producing surplus of that good and trading it to others, and in return perhaps specializing (emergence of trade).
- Check conservation: total produced = total consumed + spoiled each cycle (approximately, ignoring what's carried in surplus to next cycle).
- If we have original output data (for instance, the CSV the C code writes ⁵¹), compare a few key metrics for the same random seed. The original metrics include "living standard" (how many multiples of needs can be satisfied) which shows cycles ⁵² – we can replicate something similar and see if cycles occur.
- **Validation Against Reported Phenomena:** The report mentions certain macro outcomes qualitatively: *robust growth to a high level, then cycles, emergence of a common exchange medium (money), increasing transaction costs with specialization*, etc. ¹⁸. We should run our simulation under similar settings (300 agents, 31 goods, many cycles) and see if we observe:
 - Total utility (production) rising initially and then fluctuating in cycles.
 - Perhaps identify if one good becomes the dominant traded item (we can track how often each good is used in trade as a proxy for "money").
 - If results differ greatly, investigate if our logic deviated or if parameters need tuning.
 - Possibly consult the original code for any "knobs" we missed (like the original had some flags like `BASICROUNDELASTIC` that influence elasticity of needs). Ensure our model captures the important ones or is set to analogous defaults.
- **Benchmark Performance:** As a final test, measure how long one cycle takes for a given population (both with and without GPU, if applicable). This will inform if the simulation can run interactively. If one cycle per second is achievable for, say, 1000 agents, that's likely fine. If not, consider reducing detail or increasing hardware use (like running multiple threads if possible). The user's GPUs are powerful, so ideally we make use of at least one effectively.

All tests (except maybe heavy long-run ones) can be automated via a test suite. We will integrate these into our GitHub (e.g., using `pytest` and possibly set up a CI workflow to run them on pushes).

Phase 6: Deployment and Further Integration

- **Repository and Versioning:** Ensure the GitHub repo has a clear README explaining how to install and run the simulation and dashboard. Possibly use `requirements.txt` or a `Pipenv/Conda` environment file listing dependencies (Dash, NumPy, CuPy, etc.). The code should be structured for clarity (maybe a package `emergent_money` with modules `agent.py`, `simulation.py`, `ui.py`, etc.). - **Web Deployment (Future):** If we want to deploy the dashboard for remote access or demonstration, we can containerize it (e.g., a Dockerfile with the Python app and a GPU base image for CuPy). Or deploy to a cloud service. Note that using GPUs in the cloud can be costly, so if it's just for sharing results, we might run a smaller simulation on CPU in the cloud. Alternatively, we can capture simulation runs and present static results on a webpage. - **Continuous Improvement:** After initial success, we can iterate: add the more advanced features like the game strategies for agents (maybe integrate a reinforcement learning agent to test economic scenarios), or allow loading/saving simulation state, etc. The modular plan ensures that such additions (e.g., adding a new type of agent or a new UI panel) don't require rewriting the foundation again.

Throughout development, maintain close interaction between *planning* and *coding*. The user (as project lead) will continuously refine specifications, and the AI will provide implementation suggestions. This

synergy will be captured in commit history and design notes, creating a transparent modernization process.

Integration of Simulation, Visualization, and Deployment

With the core and UI built, we need to ensure they work seamlessly together and plan for eventual deployment beyond the development environment:

- **Connecting Simulation and UI:** We touched on this in UI development. The integration approach will be either:
 - *In-process integration:* The Dash app runs in the same Python process as the simulation. This is simplest – the callback functions directly manipulate the `Simulation` object. We just need to manage state carefully (e.g., if the simulation is running continuously in a thread, the UI should read snapshots atomically). We will likely implement a threading lock or similar if concurrent access becomes an issue.
 - *Client-server separation:* This would mean running the simulation loop in a separate thread or even process and having the UI query it (via HTTP or shared memory). This is more complex and likely unnecessary for a single-machine setup. It might be useful if we eventually separate the simulation engine (perhaps in C++/GPU code) from the UI (in a browser), but that's beyond initial scope.

We will document how the integration is done so that if someone wants to replace the UI (say, build a different front-end) they know how to interface with the simulation (likely through a small API we expose).

- **Possible Web Deployment:** If we want to allow remote users to interact with the simulation, one approach is to deploy the Dash app to a server (for example, using Heroku or AWS). The Dash app can be containerized. However, to leverage GPUs remotely, one would need a cloud VM with GPU which is expensive for continuous hosting. Alternatively, we can provide an option to run the simulation without GPU (or use smaller scale) for public demo, and keep GPU usage for local runs. We should design the code to detect if a GPU is available and if CuPy is installed, otherwise fall back to NumPy gracefully. This way the app can run in environments without CUDA. For instance, we might have:

```
try:
    import cupy as np
    CUPY_ENABLED = True
except ImportError:
    import numpy as np
    CUPY_ENABLED = False
```

Then use `np` in code for arrays. If run on a system with no GPU, it'll just use NumPy (slower but maybe okay for small runs).

- **Security & Stability:** If deploying a web app or if multiple users use it, consider any stability issues. Ensure that user inputs via the UI are validated (e.g., if we allow user to type a number for subsidy, handle non-numeric input). Since the “user” in this context is presumably the developer themselves or a controlled environment, this is a minor concern, but good practice if broadening access.

- **Source Control & Collaboration:** Continue using Git for any deployment-specific changes. Possibly use tags or releases for major versions (like v1.0 for the initial working port). If the project becomes open-source, we might integrate issue tracking for community feedback. If multiple contributors join, enforce code reviews (which AI can also help by suggesting improvements).
- **GitHub Integration:** We might use GitHub's integration with project boards to manage to-do items that come up during development. Also, if comfortable, exploring **Copilot's CLI or Chat** to generate tests or documentation might help. GitHub Actions for CI was mentioned – we can implement a basic CI that runs tests on push and maybe builds the Docker container to ensure reproducibility.

In conclusion, this integrated approach ensures the simulation core, the performance optimizations, and the interactive UI come together in a cohesive system. By following this structured plan and leveraging modern AI coding tools, the legacy C simulation will be transformed into a scalable, user-friendly application. We will gain maintainability and insight (through the dashboard) while preserving the rich dynamics of the original model. The use of GPUs and parallelism will position the simulation for larger experiments than were feasible in 2010, and the modular architecture will allow continuous enhancements, be it economic rules changes or AI-driven agent strategies, in the future.

Benchmarking, Testing, and Validation Strategy

As emphasized, validating the modernized simulation is crucial. We outline the strategy here, consolidating some points from earlier sections:

- **Baseline Verification:** Use the original code's outputs as a baseline. Since the original writes cycle-by-cycle data to `adamsmith.txt`, we can run it for a short scenario and extract key figures (e.g., total production, spoils, average prices) ⁵¹. Then run our Python simulation with the same initial conditions and random seed and compare those figures. Small differences might occur due to float vs int arithmetic or slight logic differences, but trends should align. If major discrepancies appear, inspect where our model diverges.
- **Regression Tests:** Every time we modify the code (especially performance tweaks or enabling GPU), rerun the test suite and a few sample scenarios to ensure we haven't altered the outcome. For instance, after switching to CuPy, confirm that results are bit-for-bit the same as with NumPy for a given seed (randomness handling must also be consistent – if using CuPy's random, seed it accordingly, or generate randomness on CPU and transfer). This gives confidence that optimization didn't change model behavior.
- **Stress Testing:** Test the upper limits of the simulation. Increase agent count (say to 1000 or 5000) and see if the simulation still runs correctly and how performance scales. This might reveal memory issues or algorithmic bottlenecks (e.g., an $O(N^2)$ part that becomes problematic). It's better to catch these early and possibly refine the algorithm (for example, if the friend-finding trade step is $O(N_{\text{friends}} * N_{\text{goods}}^2)$, that's fine for $N_{\text{friends}}=10$, $N_{\text{goods}}=31$, but if one day we increase goods to 100 or allow more friends, that might need a better approach or heuristics).
- **User Acceptance Testing (UAT):** The user (and any stakeholders) should try out the interactive dashboard once it's ready, to ensure it meets expectations. This includes checking that the UI is intuitive, controls work as expected, and the information displayed is useful. We might refine the

UI based on feedback – e.g., add a missing metric, or change how data is visualized (log scale for certain charts, etc.). Since this simulation is also a research/educational tool, ensuring the UI communicates the right insights (like showing the cycle phenomenon clearly) is part of validation.

- **Documentation and Reproducibility:** Document how to run experiments and interpret outputs. For testing purposes, one might include a Jupyter notebook that runs the simulation with various configurations and plots results, to demonstrate usage outside the interactive mode. This can serve both as a test and as a tutorial for users.

With all the above, by the end of the modernization project, we should have high confidence that the new implementation is *correct*, *efficient*, and *user-friendly*. We will have not only replicated the original model's results but also provided a platform to explore it further (with the user able to intervene and observe outcomes in real time).

Finally, the source control discipline ensures the codebase remains maintainable as it evolves. Each feature or fix is tracked, and we can always refer back to see why a change was made (which might be accompanied by AI explanation in commit messages or comments if we use those features).

In summary, this comprehensive plan lays out how to systematically convert the legacy Emergent Money simulation into a modern Python-based tool, using AI assistance to expedite development, and leveraging modern computing power for improved scalability. By adhering to this plan, we aim to preserve the rich dynamics of the original model while opening the door to new experiments and easier interaction with the simulation. The end result will be a robust agent-based economic simulator that runs efficiently on current hardware and provides deep insights through an interactive interface.

Sources:

- Linturi, Risto (2010). *Social Simulation of Networked Barter Economy with Emergent Money*, TKK Reports in Information and Computer Science. (Original model description and C implementation) 1 53
- *adamsmithfeb17B.c* – legacy C code for Emergent Money simulation (excerpts for structure and parameters) 42 43 34 .
- CuPy Developers. *CuPy: NumPy/SciPy for GPU*. (On using CuPy as a drop-in NumPy replacement for GPU acceleration) 14 .
- Plotly Dash Documentation – *Dash Framework Overview*. (Dash for building interactive Python dashboards) 54 .
- Mesa Documentation – *Mesa: Agent-Based Modeling in Python*. (Mesa's capabilities for ABM and interactive visualization) 19 24 .
- FLAME GPU Project – *GPU-accelerated agent-based simulation library*. (Example of advanced GPU ABM framework) 16 .

1 2 6 9 12 13 18 49 52 53 EmergentMoney.pdf

file://file-Sn9FZ8BP2qZPw8psQb8G2

3 4 5 7 8 10 11 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47 48 50

51 adamsmithfeb17B.c

file://file-Vxip8qw8SHXxXvPN9oNA1r

14 15 26 27 CuPy: NumPy & SciPy for GPU

<https://cupy.dev/>

16 FLAMEGPU/FLAMEGPU2: FLAME GPU 2 is a GPU ... - GitHub

<https://github.com/FLAMEGPU/FLAMEGPU2>

17 Fast Large-Scale Agent-based Simulations on NVIDIA GPUs with FLAME GPU | NVIDIA Technical Blog

<https://developer.nvidia.com/blog/fast-large-scale-agent-based-simulations-on-nvidia-gpus-with-flame-gpu/>

19 20 21 22 24 25 Introductory Tutorial — Mesa .1 documentation

https://mesa.readthedocs.io/stable/tutorials/intro_tutorial.html

23 54 Create Interactive Dashboard Using Python Dash | by Nur Yaumi

<https://medium.com/@nuryaumi10/build-interactive-dashboard-using-python-dash-8d145037123>